

Dr Simon D'Alfonso



INFO90002

Database Systems & Information Modelling

Week 06

SQL 2

- Nested/Sub queries
- DML
 - Comparison & Logic Operators, Set Operations, Multiple record INSERTs, INSERT from a table, UPDATE, DELETE, REPLACE
- DDL
 - ALTER and DROP, TRUNCATE, RENAME
- DCL
 - GRANT and REVOKE
- Views

- Select allows you to nest *sub-queries* inside the main or “outer” query
- A nested query is simply another Select query you write to produce a table of data
 - remember that all select queries return a “table”
- A common use of sub-queries is to perform tests
 - set membership, set comparisons
- Often there is an equivalent Join query
- Put the subquery inside round brackets

```
SELECT DISTINCT saleId FROM Sale
WHERE departmentid IN
    (SELECT departmentId FROM Department
     WHERE floor = 2);
```



- **IN / NOT IN**
 - is the value a member of the set returned by the Subquery?
- **ALL**
 - true if all values returned meet the condition
- **WHERE [NOT] EXISTS**
 - true if the subquery yields any [/ no] results

- auction example: Buyer, Seller, Artefact, Offer tables

ID	Name	Description
1	Vase	Old Vase
2	Knife	Old Knife
3	Pot	Old Pot

SellerID	Name	Phone
1	Abby	0233232232
2	Ben	0311111111
3	Carl	0333333333

BuyerID	Name	Phone
1	Maggie	0333333333
2	Nicole	0444444444
3	Oleg	0555555555

SellerID	ArtefactID	BuyerID	Date	Amount	Acceptance
1	1	1	2012-06-20	81223.23	N
1	1	2	2012-06-20	82223.23	N
2	2	1	2012-06-20	19.95	N
2	2	2	2012-06-20	23.00	N

Set queries

- which Artefacts *don't have* offers made on them?

```
SELECT * FROM Artefact
WHERE ID NOT IN
(SELECT ArtefactID FROM Offer);
```

ID	Name	Description
3	Pot	Old Pot

- which Buyers *haven't* made a bid for Artefact 3?

```
SELECT * FROM Buyer
WHERE BuyerID NOT IN
(SELECT BuyerID FROM Offer
WHERE ArtefactID = 3);
```

BuyerID	Name	Phone
1	Maggie	0333333333
2	Nicole	0444444444
3	Oleg	0555555555

- which Buyers *haven't* made a bid for the "Pot" Artefact?

```
SELECT * FROM Buyer
WHERE BuyerID NOT IN
(SELECT BuyerID FROM Offer
WHERE ArtefactID IN
(SELECT ID FROM Artefact
WHERE Name = "Pot"));
```

BuyerID	Name	Phone
1	Maggie	0333333333
2	Nicole	0444444444
3	Oleg	0555555555

- which Buyers have made a bid for the “Knife” Artefact?

```
SELECT * FROM Buyer
  WHERE BuyerID IN
    (SELECT BuyerID FROM Offer
      WHERE ArtefactID IN
        (SELECT ID FROM Artefact
          WHERE Name = "Knife"));
```

BuyerID	Name	Phone
1	Maggie	0333333333
2	Nicole	0444444444

There is often an equivalent Join that will achieve the same result. The above is equivalent to:

```
SELECT Buyer.*
FROM Buyer NATURAL JOIN Offer NATURAL JOIN Artefact
WHERE Artefact.name = 'Knife' ;
```

These functions operate on a set of values (e.g. in a column of a table) and return a single value

- AVG()
 - Average value
- MIN()
 - Minimum value
- MAX()
 - Maximum value
- COUNT()
 - Number of values
- SUM()
 - Sum of values

- and there are others ...
 - <http://dev.mysql.com/doc/refman/5.7/en/group-by-functions.html>
- These ignore null values, and return null if all values are null.
- But COUNT(*) counts the rows not the values, and thus even if the value is NULL it is still counted.

Finding above average

- Consider the Item table in our labs database
- Which items have a price that is higher than the average?

Item	
ItemID	SMALLINT
Name	VARCHAR(50)
Type	CHAR(1)
Colour	VARCHAR(20)
ItemPrice	DECIMAL(9,2)

```
SELECT * FROM Item
WHERE itemPrice >
      (SELECT AVG(itemPrice) FROM Item);
```

234.766400

itemID	Name	Type	Colour	itemPrice
1	Boots Riding	C	Brown	235.00
2	Horse saddle	R	Brown	1895.00
12	Gortex Rain Coat	C	Green	249.75
19	Tent - 2 person	F	Khaki	399.95
20	Tent - 8 person	F	Khaki	785.96
21	Tent - 4 person	F	Blue	638.95
24	Boots - Womens Goretex	C	Grey	289.95
25	Boots - Mens Hiking	C	Grey	299.95

Finding a maximum

- Which item has the highest cost?

```
SELECT * FROM Item
WHERE itemPrice =
    (SELECT MAX(itemPrice) FROM Item);
```

1,895

Item	
ItemID	SMALLINT
Name	VARCHAR(50)
Type	CHAR(1)
Colour	VARCHAR(20)
ItemPrice	DECIMAL(9,2)

```
SELECT * FROM Item
ORDER BY itemprice DESC
LIMIT 1;
```

itemID	Name	Type	Colour	itemPrice
2	Horse saddle	R	Brown	1895.00

- Will these two methods always give the same answer?

- another method

```
SELECT * FROM Item
WHERE itemPrice >= ALL
      (SELECT itemPrice FROM Item);
```

- and another: a “correlated subquery”

```
SELECT * FROM Item A
WHERE itemPrice > ALL
      (SELECT itemPrice FROM Item B
       WHERE A.itemId != B.itemId);
```

- SQL keywords are *not* case-sensitive.
 - the traditional convention is to CAPITALISE them for clarity
- Table names *are* case sensitive in Unix, but not Windows (and possibly *not* case-sensitive if you use the InnoDB storage engine)
 - Account <> account <> ACCOUNT (in Unix)
- Column names are *not* case-sensitive
 - ACCOUNTID == AccountID == AcCoUnTID
- Case-sensitivity of DATA ('strings in quotes') depends on character set used.
(The default 'latin1' set is *not* case-sensitive.)
- SQL handles expressions including maths:
 - SELECT 1*2+3/4-5;
 - SELECT now();

- Comparison

Operator	Description
=	Equal to
<	Less than
>	Greater than
<=	Less than or equal to
>=	Greater than or equal to
<> OR !=	Not equal to

- Logic

- SQL supports AND, NOT, OR logical operators

- `SELECT * FROM Furniture`
`WHERE ((Type= 'Chair' AND Colour = 'Black')`
`OR NOT (Type = 'Lamp' AND Colour = 'White'));`

- We can combine results from two or more queries that return *the same number of columns* - although it usually only makes sense if they are the *same columns*.
- UNION
 - Show all rows returned from the queries, without duplicates
- INTERSECT
 - Show only rows that are common in the queries
- EXCEPT
 - Show only rows that are different in the queries
- [UNION/INTERSECT/EXCEPT] ALL
 - If you want duplicate rows shown in the results you need to use the ALL keyword, e.g. UNION ALL.
- In MySQL only UNION and UNION ALL are supported

UNION example

```
SELECT * FROM Department
WHERE floor = 1
UNION
SELECT * FROM Department
WHERE floor = 3;
```

DepartmentID	DepartmentName	DepartmentFloor	DepartmentPhone	ManagerID
6	Navigation	1	41	3
8	Books	1	81	4
4	Equipment	3	57	3
NULL	NULL	NULL	NULL	NULL

(what if the subsets overlap?)

- **FORMAT()**
 - changes format of output of Select
 - e.g. **FORMAT (N, D)**
 - N: A number which may be an integer, a decimal or a float.
 - D: How many decimals the output contains
 - **FORMAT(123456.1234, 2)** gives '123,456.12'
- **CAST()**
 - changes data type of output
 - e.g. **CAST (Expression AS Type)**
 - **CAST("1234.55" AS UNSIGNED)** Gives 1235
 - **CAST("1234.55" AS DECIMAL(7,1))** Gives 1234.6
 - Valid types include
 - **BINARY[(N)], CHAR[(N)], DATE, DATETIME, DECIMAL[(M[,D])], SIGNED, TIME, UNSIGNED**

Output without Format

```
SELECT Department.DepartmentID, SUM(EmployeeSalary*Bonus) AS TotSalary
FROM Department INNER JOIN Employee ON Department.DepartmentID = Employee.DepartmentID
GROUP BY Department.DepartmentID;
```

DepartmentID	TotSalary
1	67499.9982118607
2	60000
3	32639.9993896484
4	27039.9990081787
5	15000
6	15000
7	16500.0003576279
8	15149.9998569489
9	99000
10	35000
11	101200.002193451

messy



```
SELECT Department.DepartmentID, FORMAT(SUM(EmployeeSalary*Bonus),2) AS TotSalary  
FROM Department INNER JOIN Employee ON Department.DepartmentID = Employee.DepartmentID  
GROUP BY Department.DepartmentID;
```

but Format() converts
numbers to strings ...

what happens now
if we sort by TotSalary?

DepartmentID	TotSalary
1	67,500.00
2	60,000.00
3	32,640.00
4	27,040.00
5	15,000.00
6	15,000.00
7	16,500.00
8	15,150.00
9	99,000.00
10	35,000.00
11	101,200.00

Formatting output (Format)

```
SELECT Department.DepartmentID, FORMAT(SUM(EmployeeSalary*Bonus),2) AS TotSalary
FROM Department INNER JOIN Employee ON Department.DepartmentID = Employee.DepartmentID
GROUP BY Department.DepartmentID
ORDER BY TotSalary DESC;
```

DepartmentID	TotSalary
9	99,000.00
1	67,500.00
2	60,000.00
10	35,000.00
3	32,640.00
4	27,040.00
7	16,500.00
8	15,150.00
6	15,000.00
5	15,000.00
11	101,200.00

wrong



Formatting output (Cast)

```
SELECT Department.DepartmentID, CAST(SUM(EmployeeSalary*Bonus) AS DECIMAL(9,2)) AS TotSalary
FROM Department INNER JOIN Employee ON Department.DepartmentID = Employee.DepartmentID
GROUP BY Department.DepartmentID
ORDER BY TotSalary DESC;
```

DepartmentID	TotSalary
11	101200.00
9	99000.00
1	67500.00
2	60000.00
10	35000.00
3	32640.00
4	27040.00
7	16500.00
8	15150.00
6	15000.00
5	15000.00

These are numbers, so
ordering works again

- IFNULL()
 - Can convert a null to a zero (can be useful in calculations)
 - `SELECT 1 + IFNULL(wagevalue, 0)`
 - gives 1+0 for null fields, and 1+wagevalue for non null fields
 - failure to do this results in a NULL answer for values where wagevalue is NULL

(example on next two slides)

```

SELECT e.ID, e.Name, e.Address, DateHired, DateLeft,
       EmployeeType, ContractNumber, BillingRate,
       AnnualSalary, StockOption, HourlyRate
FROM Employee e
LEFT OUTER JOIN Hourly h ON e.ID = h.ID
LEFT OUTER JOIN Salaried s ON e.ID = s.ID
LEFT OUTER JOIN Consultant c ON e.ID = c.ID;

```

ID	Name	Address	DateHired	DateLeft	EmployeeType	ContractNumber	BillingRate	AnnualSalary	StockOption	HourlyRate
1	Sean	Sean's Address	2012-02-02	NULL	S	NULL	NULL	92000.00	N	NULL
2	Linda	Linda's Address	2011-06-12	NULL	S	NULL	NULL	92300.00	Y	NULL
3	Alice	Alice's Address	2012-12-02	NULL	H	NULL	NULL	NULL	NULL	23.43
4	Alan	Alan's Address	2010-01-22	NULL	H	NULL	NULL	NULL	NULL	29.43
5	Peter	Peter's Address	2010-09-07	NULL	C	19223	210.00	NULL	NULL	NULL
6	Rich	Rich's Address	2012-05-19	NULL	C	19220	420.00	NULL	NULL	NULL

```

SELECT e.ID, e.Name, e.Address, DateHired,
       EmployeeType, IFNULL(ContractNumber,0) ContractNbr,
       IFNULL(BillingRate,0) BillRate, IFNULL(AnnualSalary,0) Salary,
       IFNULL(StockOption,"") StockOpt, IFNULL(HourlyRate,0) HrlyRate
FROM Employee e
LEFT OUTER JOIN Hourly h ON e.ID = h.ID
LEFT OUTER JOIN Salaried s ON e.ID = s.ID
LEFT OUTER JOIN Consultant c ON e.ID = c.ID;

```

	ID	Name	Address	DateHired	EmployeeType	ContractNbr	BillRate	Salary	StockOpt	HrlyRate
▶	1	Sean	Sean's Address	2012-02-02	S	0	0.00	92000.00	N	0.00
	2	Linda	Linda's Address	2011-06-12	S	0	0.00	92300.00	Y	0.00
	3	Alice	Alice's Address	2012-12-02	H	0	0.00	0.00		23.43
	4	Alan	Alan's Address	2010-01-22	H	0	0.00	0.00		29.43
	5	Peter	Peter's Address	2010-09-07	C	19223	210.00	0.00		0.00
	6	Rich	Rich's Address	2012-05-19	C	19220	420.00	0.00		0.00

- LOWER() / UPPER()
 - Change string to lower / upper case
 - e.g. `SELECT LOWER('That')` gives 'that'
 - `SELECT UPPER('That')` gives 'THAT'
- LEFT() / RIGHT()
 - Returns the leftmost / rightmost N characters from a string
 - e.g. `SELECT LEFT('This is a test', 6)` gives "This i"
 - e.g. `SELECT RIGHT('This is a test', 6)` gives "a test"
- Date and time functions
 - <http://dev.mysql.com/doc/refman/5.5/en/date-and-time-functions.html>
 - including `DATEDIFF()`, `TIMEDIFF()`, `NOW()` or `TIMESTAMP()`, `CURDATE()`, `CURTIME()`

- Inserting records from another table
 - Note: table must already exist

```
INSERT INTO NewEmployee  
SELECT * FROM Employee;
```

- Insert multiple rows

```
INSERT INTO Employee VALUES  
(DEFAULT, "A", "A's Addr", "2012-02-02", NULL, "S"),  
(DEFAULT, "B", "B's Addr", "2012-02-02", NULL, "S"),  
(DEFAULT, "C", "C's Addr", "2012-02-02", NULL, "S");
```

```
INSERT INTO Employee  
(Name, Address, DateHired, EmployeeType)  
VALUES  
("D", "D's Addr", "2012-02-02", "C"),  
("E", "E's Addr", "2012-02-02", "C"),  
("F", "F's Addr", "2012-02-02", "C");
```

- Be careful to specify a WHERE clause
 - unless you want it to operate on EVERY row in the table

```
UPDATE Hourly
    SET HourlyRate = HourlyRate * 1.10;
```

- Increase salaries greater than \$100k by 10% and all other salaries by 5%

```
UPDATE Salaried
    SET AnnualSalary = AnnualSalary * 1.05
    WHERE AnnualSalary <= 100000;
UPDATE Salaried
    SET AnnualSalary = AnnualSalary * 1.10
    WHERE AnnualSalary > 100000;
```

- Any problems with this?

- A better solution is to use the CASE expression

```
UPDATE Salaried
  SET AnnualSalary =
    CASE
      WHEN AnnualSalary <= 100000
      THEN AnnualSalary * 1.05
      ELSE AnnualSalary * 1.10
    END;
```

- now we process each row independently, one at a time

- CASE can also be used in SELECT statements
- e.g “Calculate our annual bonuses. Give each employee a 10% bonus, except those who work in Clothes or Books, who get 20%.”

```
1 • SELECT employeeId, lastName, departmentId, salary,
2 CASE
3     WHEN departmentId in
4         (SELECT departmentId FROM Department WHERE name in ('clothes', 'books'))
5     THEN salary * 0.2
6     ELSE salary * 0.1
7 END as bonus
8 FROM employee
9 ORDER BY departmentid;
```

Result Grid

	employeeId	lastName	departmentId	salary	bonus
▶	1	Munro	1	125000.00	12500.000
	11	Skeeter	2	45000.00	9000.000
	13	Smith	3	46000.00	9200.000
	12	Montez	3	46000.00	9200.000
	15	Mason	4	45000.00	4500.000
	14	Innit	4	41000.00	4100.000

Result 11 x

Yes/No Questions using CASE

- You can use CASE to answer yes/no or true/false questions.
- e.g “Are there more than ten customers?”

```

1  /* Are there more than ten customers? */
2  • SELECT
3  CASE
4      WHEN COUNT(*) > 10
5      THEN 'yes'
6      ELSE 'no'
7  END as answer
8  FROM Customer;

```

Result Grid

answer
no

```

1  /* General true/false question */
2  • SELECT
3  CASE
4      WHEN 1 = 2
5      THEN 'true'
6      ELSE 'false'
7  END as answer
8  ;

```

Result Grid

answer
false

- REPLACE
 - REPLACE works the same as INSERT
 - EXCEPT that if an old row in a table has a key value the same as the new row, then it is overwritten...
- DELETE
 - be careful to use a WHERE clause ... What does this do?

```
DELETE FROM Employee;
```

 - Usually you should do use a filter:

```
DELETE FROM Employee  
WHERE Name = "Grace";
```
 - If you delete a row that has rows in other tables dependent on it, either:
 - the dependent rows are deleted too, or
 - the dependent rows get 'null' or a default, or
 - your attempt to delete is blocked
 - you decide what action to take when you set up the tables
 - ON DELETE CASCADE or ON DELETE RESTRICT...

- a View is a select statement that persists, and can be treated as though it were a table by other SQL statements
- Used to:
 - hide the complexity of queries from users
 - hide structure of data from users
 - hide data from users
 - different users use different views
 - e.g. allow someone to access employee table, but not salaries column
 - one way of improving database security
- To create a view...
 - **CREATE VIEW** nameofview **AS** validSelectStatement
 - its definition (but not its output) is stored in the database
 - can be used as though it is a table

```
CREATE VIEW DepartmentSales AS
SELECT departmentId, name, COUNT(*) as numSales
FROM Department NATURAL JOIN Sale
GROUP BY departmentId;
```

6 • SELECT * FROM DepartmentSales;

Result Grid			
Filter Rows:			
Export: Wrap Cell Content:			
	departmentId	name	numSales
▶	2	Books	6
	3	Clothes	8
	4	Equipment	6
	5	Furniture	4
	6	Navigation	13
	7	Recreation	6

6 • SELECT * FROM DepartmentSales
7 WHERE numSales > 5;


Result Grid			
Filter Rows:			
Export: Wrap Cell Content:			
	departmentId	name	numSales
▶	2	Books	6
	3	Clothes	8
	4	Equipment	6
	6	Navigation	13
	7	Recreation	6

- Conditions that must be satisfied:
 - the select clause only contains attribute names
 - not expressions, aggregates or distinct
 - any attributes not listed in the select clause can be set to null
 - the query does not have a group by or having clause
- MySQL conditions for updatable views are quite stringent
 - see <http://dev.mysql.com/doc/refman/5.0/en/view-updatability.html>

Updating a View: example

Underlying
base table

->



Id	Name	City	Salary
1	John Lennon	Sydney	100000
2	Paul McCartney	Melbourne	80000
3	George Harrison	Melbourne	90000
4	Ringo Starr	Brisbane	110000

CREATE VIEW MelbRestricted AS
(SELECT id, name, city from Employee
WHERE city = 'Melbourne');

SELECT * FROM MelbRestricted;

Id	Name	City
2	Paul McCartney	Melbourne
3	George Harrison	Melbourne

INSERT INTO MelbRestricted VALUES
(null, 'Yoko Ono', 'Melbourne');

Id	Name	City	Salary
1	John Lennon	Sydney	100000
2	Paul McCartney	Melbourne	80000
3	George Harrison	Melbourne	90000
4	Ringo Starr	Brisbane	110000
6	Yoko Ono	Melbourne	NULL

- (beyond CREATE)
- ALTER
 - Allows us to add or remove columns from a table
 - **ALTER TABLE** TableName **ADD** AttributeName AttributeType
 - **ALTER TABLE** TableName **DROP** AttributeName
 - not supported by all vendors (MySQL supports it)
- RENAME
 - Allows the renaming of tables
 - **RENAME TABLE** CurrentTableName **TO** NewTableName

- TRUNCATE
 - like “DELETE FROM table” but it does more
 - differences are vendor-specific, see <http://stackoverflow.com/questions/139630/whats-the-difference-between-truncate-and-delete-in-sql> and <https://dev.mysql.com/doc/refman/5.0/en/truncate-table.html>
 - in MySQL, resets auto_increment PKs
 - cannot ROLL BACK a TRUNCATE command
 - have to get data back from backup...
- DROP
 - potentially DANGEROUS
 - Removes the table definition and the data in the table
 - There is NO UNDO COMMAND! (have to restore from backup)
 - DROP TABLE TableName

- DCL
 - Users and permissions
 - **CREATE USER, DROP USER**
 - **GRANT, REVOKE**
 - **SET PASSWORD**
- Other commands offered
 - Database administration
 - **BACKUP TABLE, RESTORE TABLE**
 - **ANALYZE TABLE**
 - Miscellaneous
 - **DESCRIBE tablename**
 - **USE db_name**
 - MySql calls these
'Database Administration Statements'

- Data Definition Language (DDL)
 - To define and set up the database
 - CREATE, ALTER, DROP
 - Also TRUNCATE, RENAME
- Data Manipulation Language (DML)
 - To maintain and use the database
 - SELECT, INSERT, DELETE, UPDATE
 - MySQL also provides others.... eg REPLACE
- Data Control Language (DCL)
 - To control access to the database
 - GRANT, REVOKE
- Other Commands
 - Administer the database
 - Transaction Control